

LOFT: A Lock-free and Adaptive Learned Index with High Scalability for Dynamic Workloads

Yuxuan Mo, Yu Hua

Huazhong University of Science and Technology

20th European Conference on Computer Systems (EuroSys), 2025

Dynamic Workloads

- Contain insert operations
 - *Growth* in the data size
 - *Changes* in data distribution
- Widely exist in real-world applications
 - e.g., Facebook, Twitter, etc.
 - Some are write-heavy^[1]




[1] Dynamo: amazon's highly available key-value store, SOSP'07

Memory Systems

- Memory systems play a **critical role** in compute systems

- High-speed CPUs
- Low-speed storage systems




- In-memory index structures contribute to overall performance

- e.g., B⁺-tree and hash maps
- Efficient data management with fast query performance




Dilemma: Data Growth vs DRAM Scaling

- **Rapid growth** of stored data^[1]



- **Slowdown** of DRAM scaling technology^[2]



[1] <https://www.statista.com/statistics/871513/worldwide-data-created/>


[2] <https://jcmit.net/memoryprice.htm>

Demand: Space-efficient and Scalable Index Structures

- Tree-like indexes



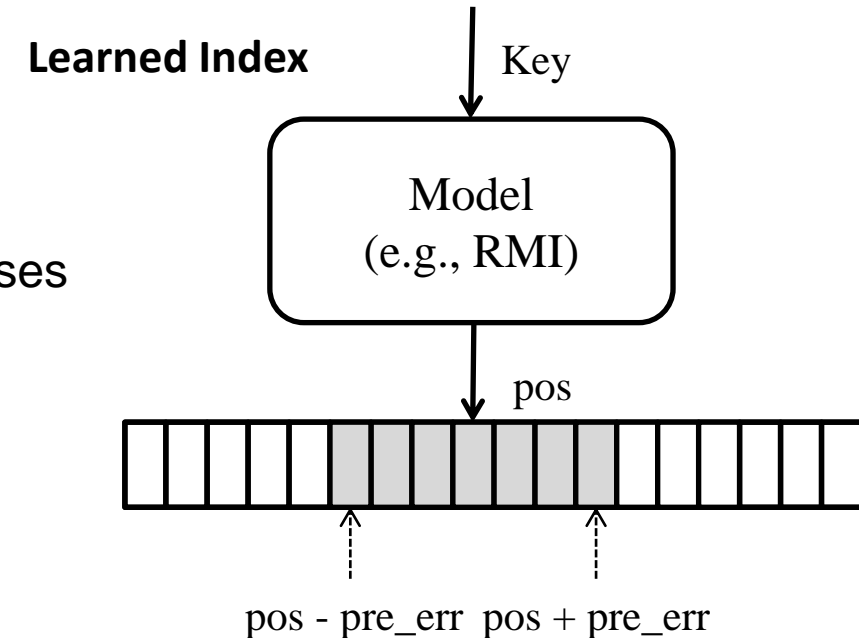
- Multiple pointer chasing operations
- Large space overhead



- Learned indexes



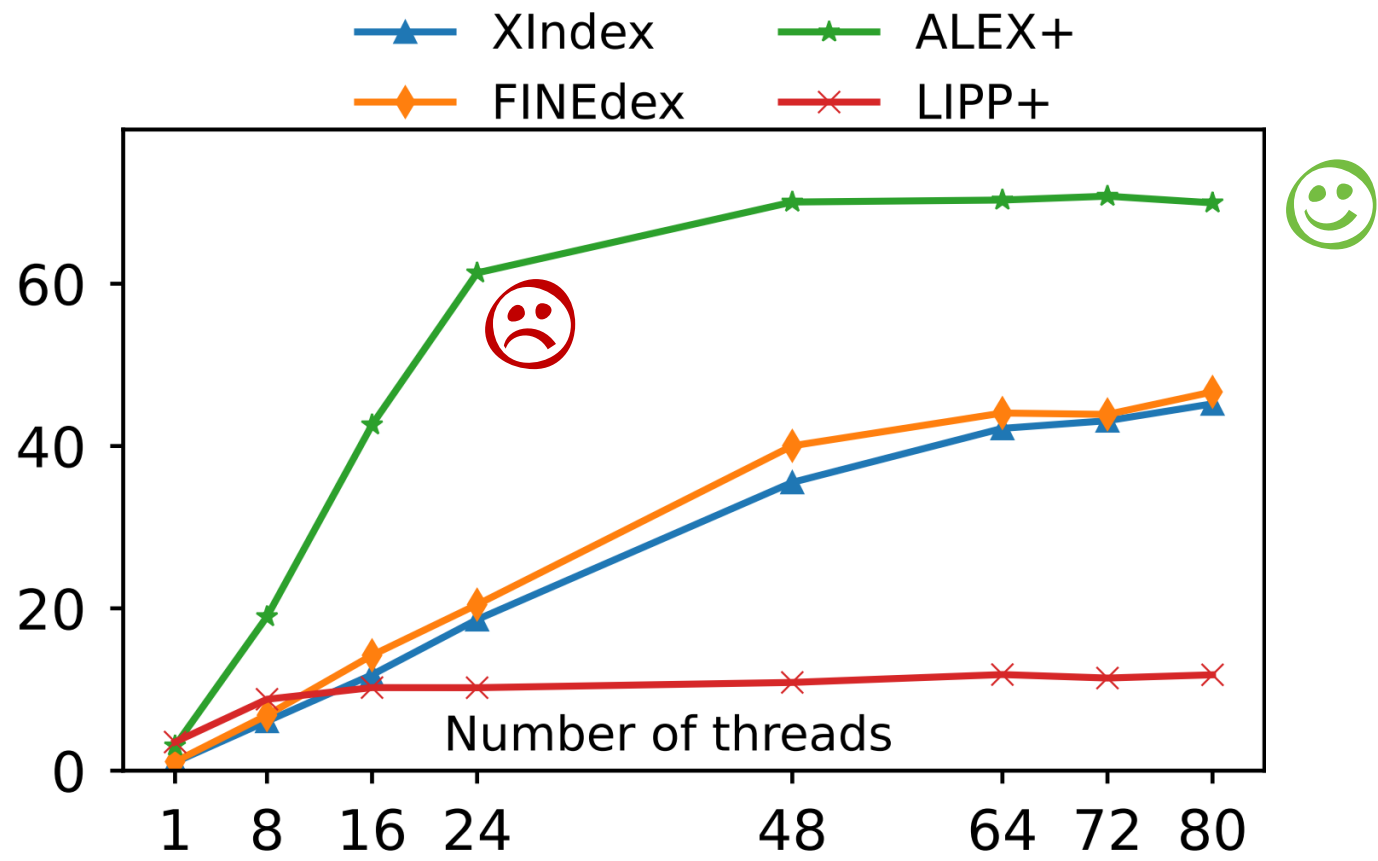
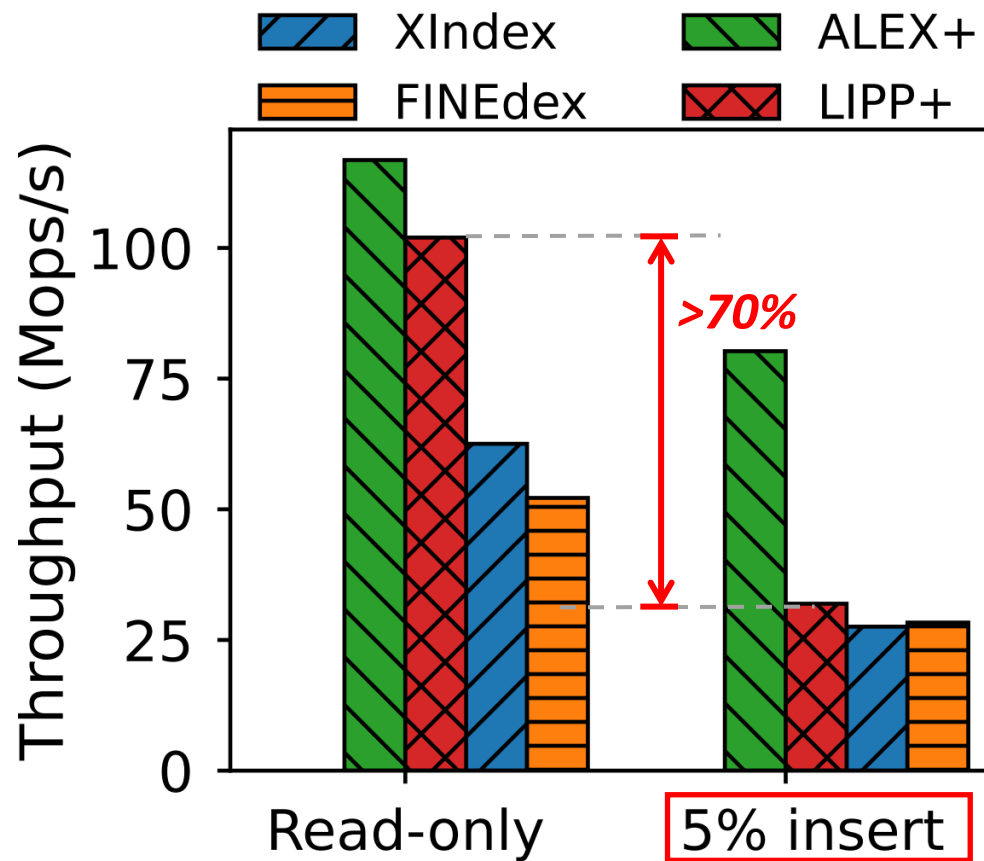
- Model-based calculation
- Computation memory



Is the learned index the optimal solution?

Existing Learned Indexes in Dynamic Workloads

- Fail to scale to dynamic workloads
- Fail to simultaneously achieve high throughput and high scalability



- The used workloads generated from YCSB.

- XIndex@PPoPP'20, FINEdex@VLDB'21, ALEX+@VLDB'22, LIPP+@VLDB'22

Existing Learned Indexes in Dynamic Workloads

- ✓ ➤ Space-efficient
 - Small number of parameters
- ✓ ➤ Efficient query performance
 - Model-based calculation
- ✗ ➤ High performance in dynamic workloads
 - High throughputs
 - High scalability



Challenge 1: Interference from Insertions

• In-place Insertion




- Model-based insertion
- **Good** query performance
- Lock-based design(ALEX+@VLDB'22)
 - **Poor** scalability

• Out-of-place Insertion



- Buffer-based insertion
- **Good** scalability
- Buffer-based design(XIndex@PPOPP'20)
 - **Poor** query performance



Existing schemes interfere with concurrent or subsequent reads.

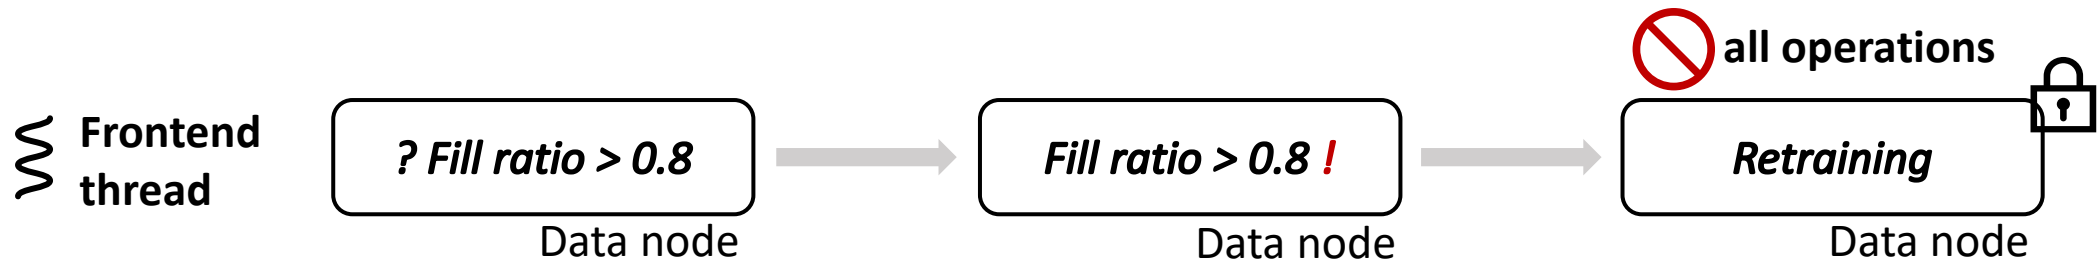
Challenge 2: Collisions between Indexing and Retraining

- Blocking retraining scheme (ALEX+@VLDB'22)

- Actively trigger retraining once the condition is met
- Block the following operations to the retrained node
- In the critical path

timely

Long tail latency



Challenge 2: Collision between Indexing and Retraining

- Blocking retraining scheme (ALEX+@VLDB'22)

- Actively trigger retraining once the condition is met

Timely

- Block the following operations to the retrained node
- In the critical path

Long tail latency


- Non-blocking retraining scheme (XIndex@PPOPP'20)

- Periodically check the data nodes using background threads

Non-blocking

- Unable to handle heavy tasks in write-intensive workloads

Long average latency



How to achieve in-time and lightweight retraining?

Challenge 3: Fixed Parameters vs Diverse Access Patterns

Write-intensive



- **Static triggering mechanism:**

- Perform retraining once the predefined condition is met

Trigger retraining in advance?

Write-intensive



- **Static retraining parameters**

- Use fixed parameters based on preliminary experiments

Preserve more free slots in advance?

Our Solution: **LOFT**

➤ To achieve high performance in **dynamic workloads**:

- **C1**: Interference introduced by insertions

Error-bounded insertion

- **C2**: Collisions between indexing and retraining

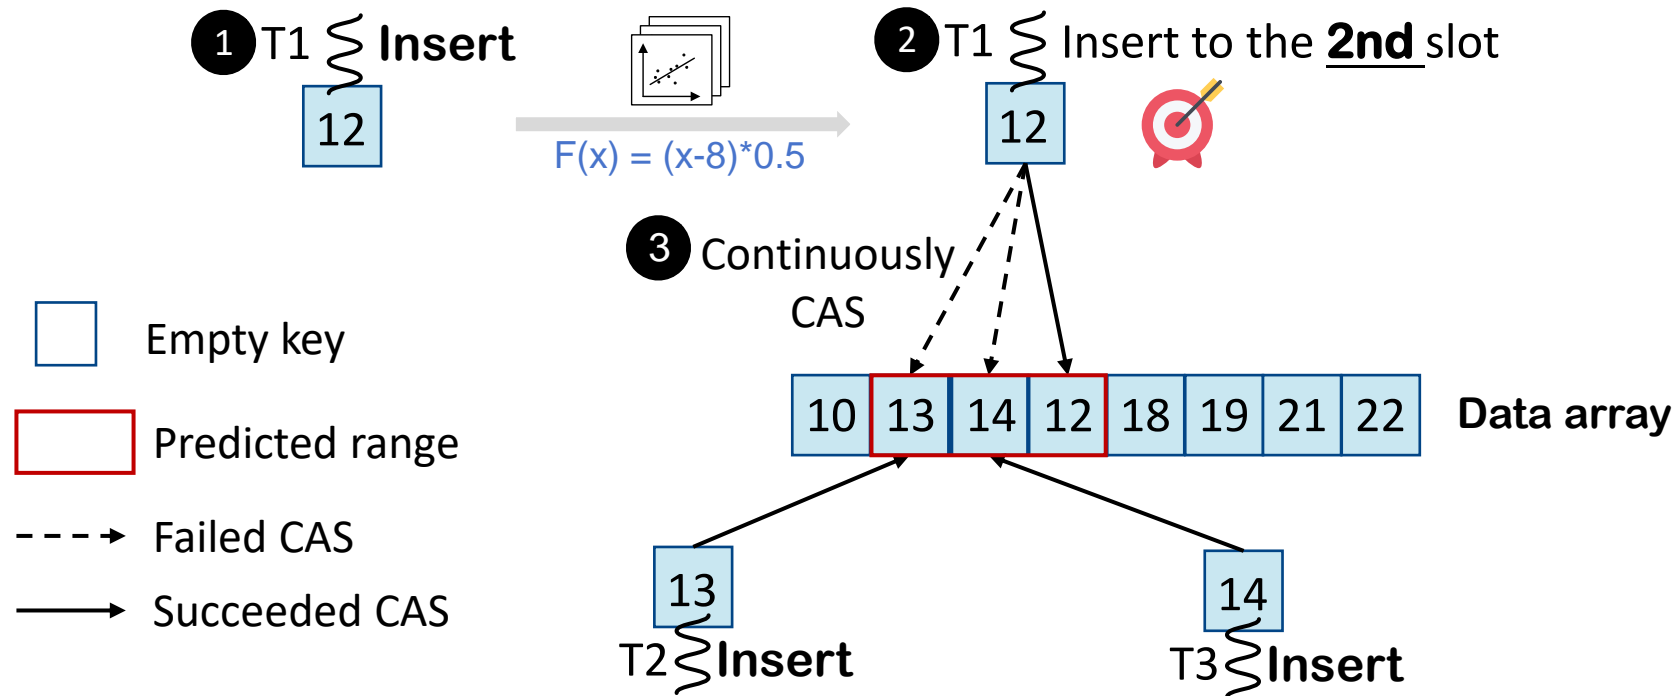
Lock-free retraining

- **C3**: Mismatch between fixed parameters and diverse access patterns

Self-tuning retraining

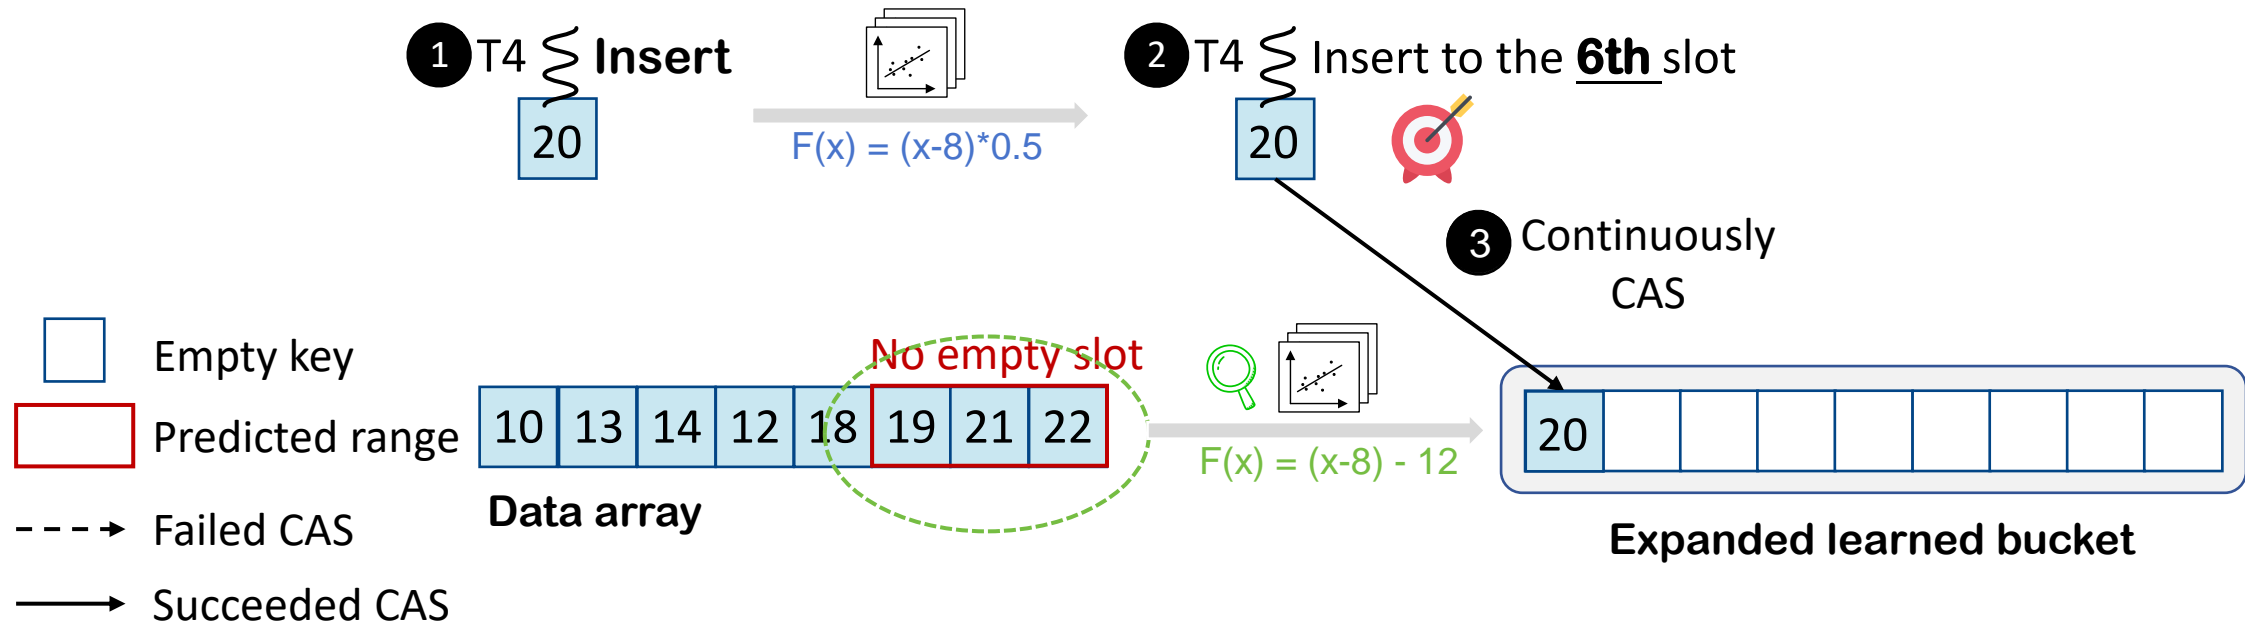
LOFT: Error-bounded Insertion

- Using **CAS*** to compete for an empty slot within the predicted range
 - No **shifting** for sorting
 - No duplicate keys



LOFT: Error-bounded Insertion

- **Expanded Learned Bucket** for possible overflows
 - A small data array with expanded models
 - Increase the expansion factor as the bucket level rises



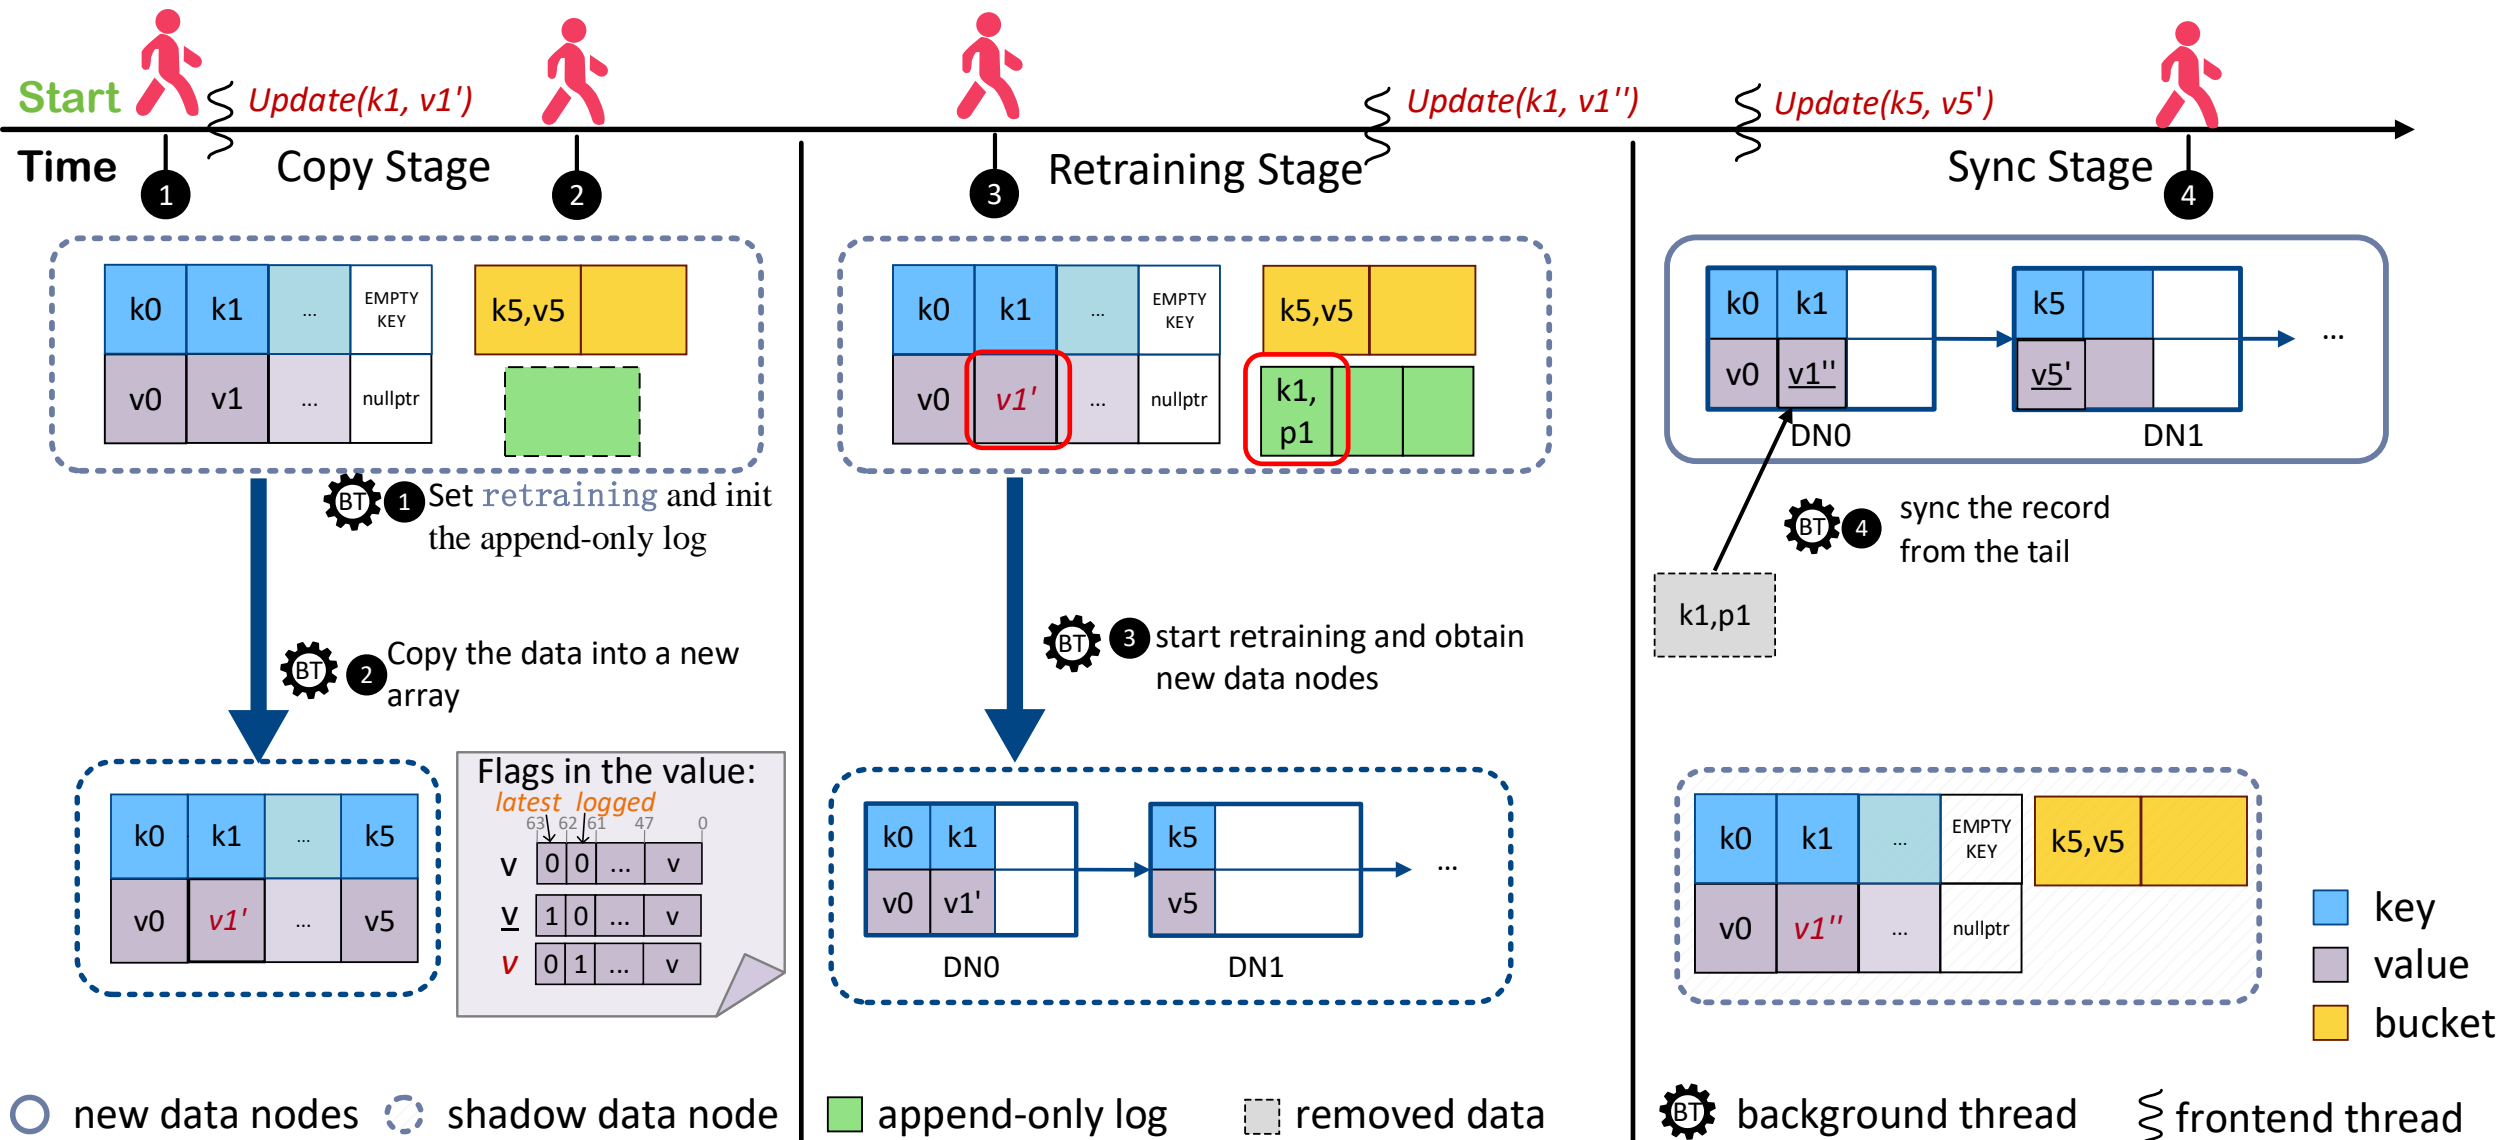
Concurrent insert operations are executed in a lock-free manner.

LOFT: Lock-free Index Operations

- Decrease read performance to minimize operation interference
- Read
 - Linear search within the predicted range
 - Reasonable overheads
- In-place update
 - Atomically update the 8-byte value pointers
- Soft delete
 - Maintain the key in the data array
 - Invalidate the value

All index operations are executed in a lock-free manner.

LOFT: Non-blocking Retraining Process



LOFT: Self-tuning Retraining

- **Write-intensive**

- Increase the expansion factor of data nodes
- Increase the predicted range

↘ Retraining frequency

- **Cold**

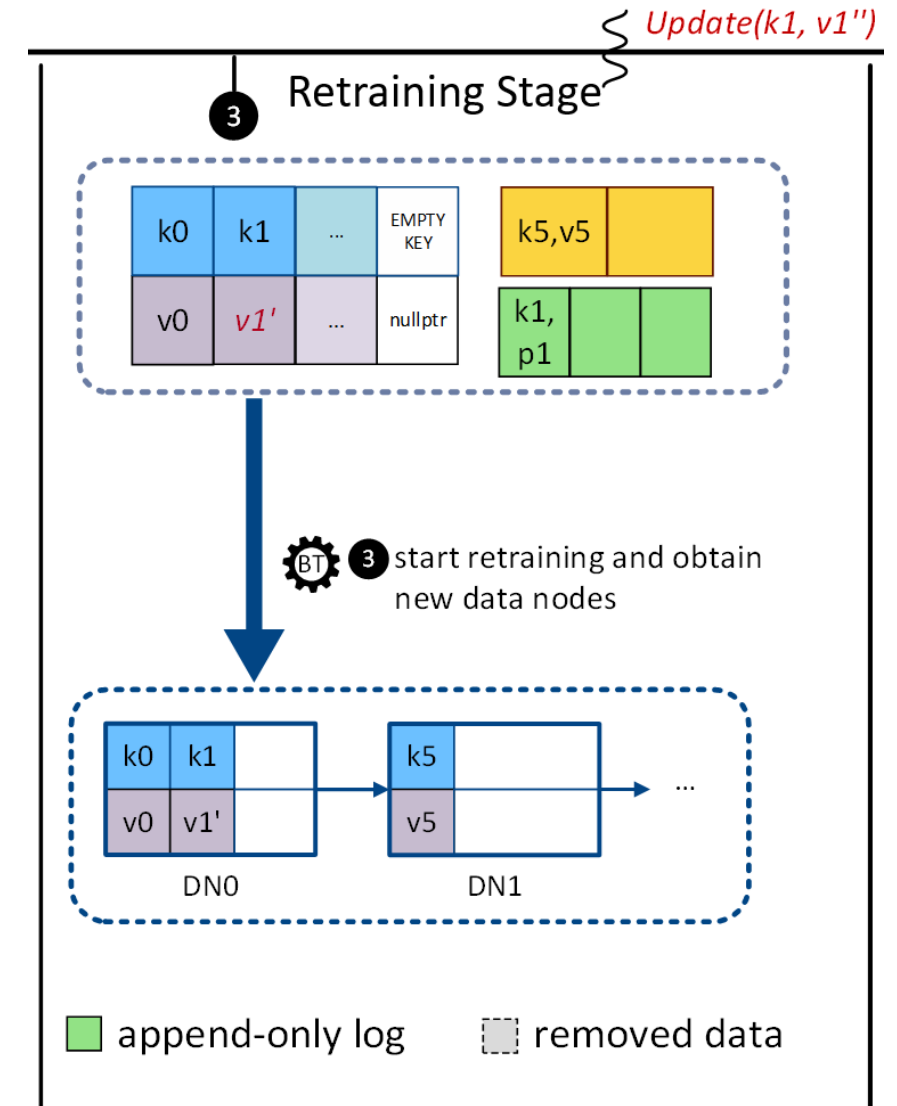
- Decrease the expansion factor
- Increase the predicted range

↘ Index size

- **Read-intensive**

- Decrease the predicted range

↘ Search length



More Details about LOFT

- Concurrency correctness
- Structure modification operations
- Informed decision making
-



Experimental Setup

- **Testbeds**

- Two 26-core Intel(R) Xeon(R) CPU @2.10GHz
- Assign one background thread to every twelve worker threads

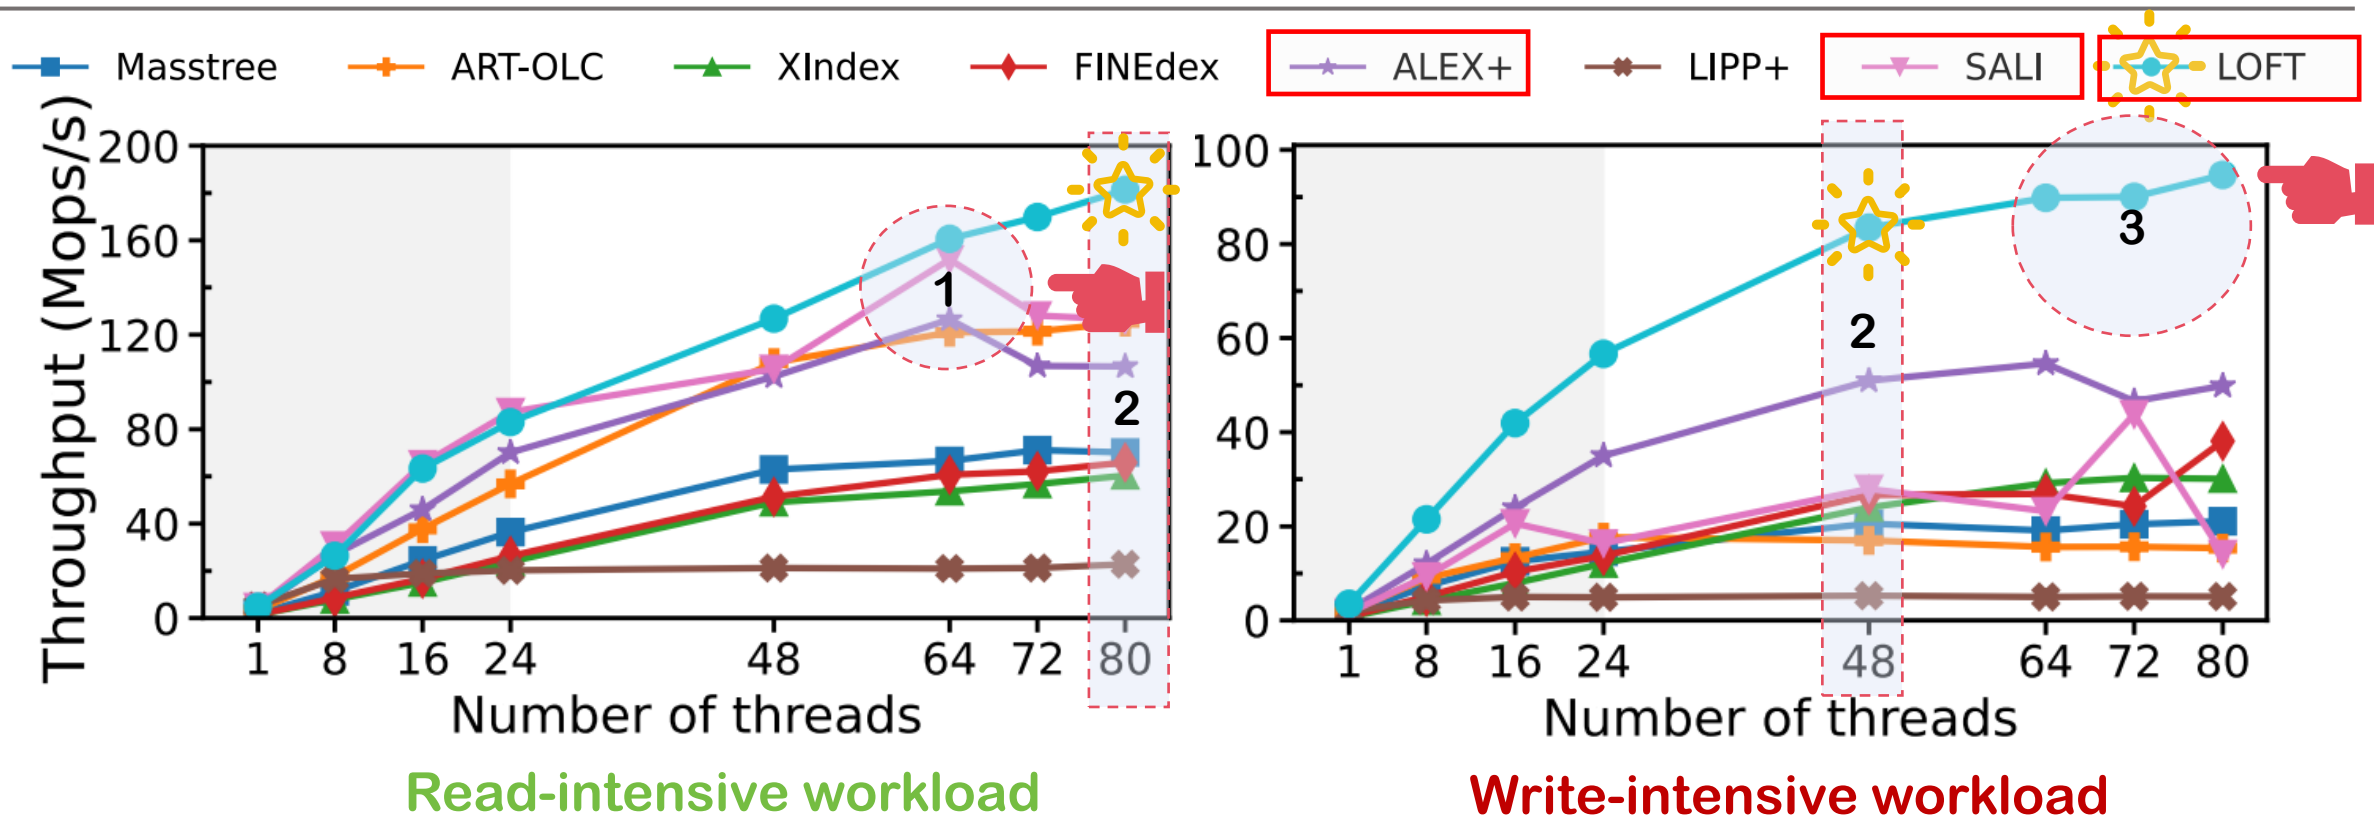
- **Workloads**

- YCSB with Zipfian distribution
- Multiple real-world datasets

- **Comparisons**

- *Conventional*: Masstree [Eurosys'12], ART-OLC [DaMoN '16]
- *Learned*: DyTIS [Eurosys'23], XIndex [PPOPP'20], FINEdex [VLDB'21], ALEX+ [VLDB'22], LIPP+ [VLDB'22], SALI [SIGMOD'23]

Evaluation on Scalability




1. Due to the in-place insertion design, ALEX+, SALI and LOFT achieve higher throughput.

2. Due to the lock-free design, LOFT achieves the best scalability.

3. LOFT improves the throughput by 1.7x – 14x on average.

Evaluation on Adaptiveness



1. The average throughputs of all indexes decline as the proportion of insertions increases.

2. Our lock-free retraining scheme enables LOFT to avoid severe performance jitter.

3. LOFT illustrates long-term stability thanks to self-tuning retraining mechanism.

Summary

- Existing learned indexes show limited scalability in dynamic workloads.
 - Display sharp performance degradation
 - Fail to simultaneously achieve high throughput and high scalability
- **LOFT**: a Lock-free and scalable learned index.
 - Error-bounded insertion scheme
 - Lock-free index operations and retraining process
 - Self-tuning retraining mechanism
- **LOFT** significantly improves the throughput with high scalability compared with state-of-the-art schemes.

Thanks!

Q & A



yuxuanmo@hust.edu.cn