

Semi-hierarchical Semantic-aware Storage Architecture

Yu Hua

Huazhong University of Science
and Technology

<http://cs.hust.edu.cn/stlab/csyhua/>

Current and Future Storage

S M I L E



Current and Future Storage 1

- SMILE
- Scale : Big Data , Big Storage

Current and Future Storage 2

- SMILE
- NN(M)-Intelligent :

Current and Future Storage 3

- SMILE
- Integrated :
 - Near Data Processing :
 - Processing in-memory (PIM)
 - In-storage computing (ISC)
 - Quantx(Micron), Optane(Intel), NDP(HUAWEI),

Current and Future Storage 4

- SMILE
- Long-term :
 - Storage media and runtime context
 - Time-sensitive and value

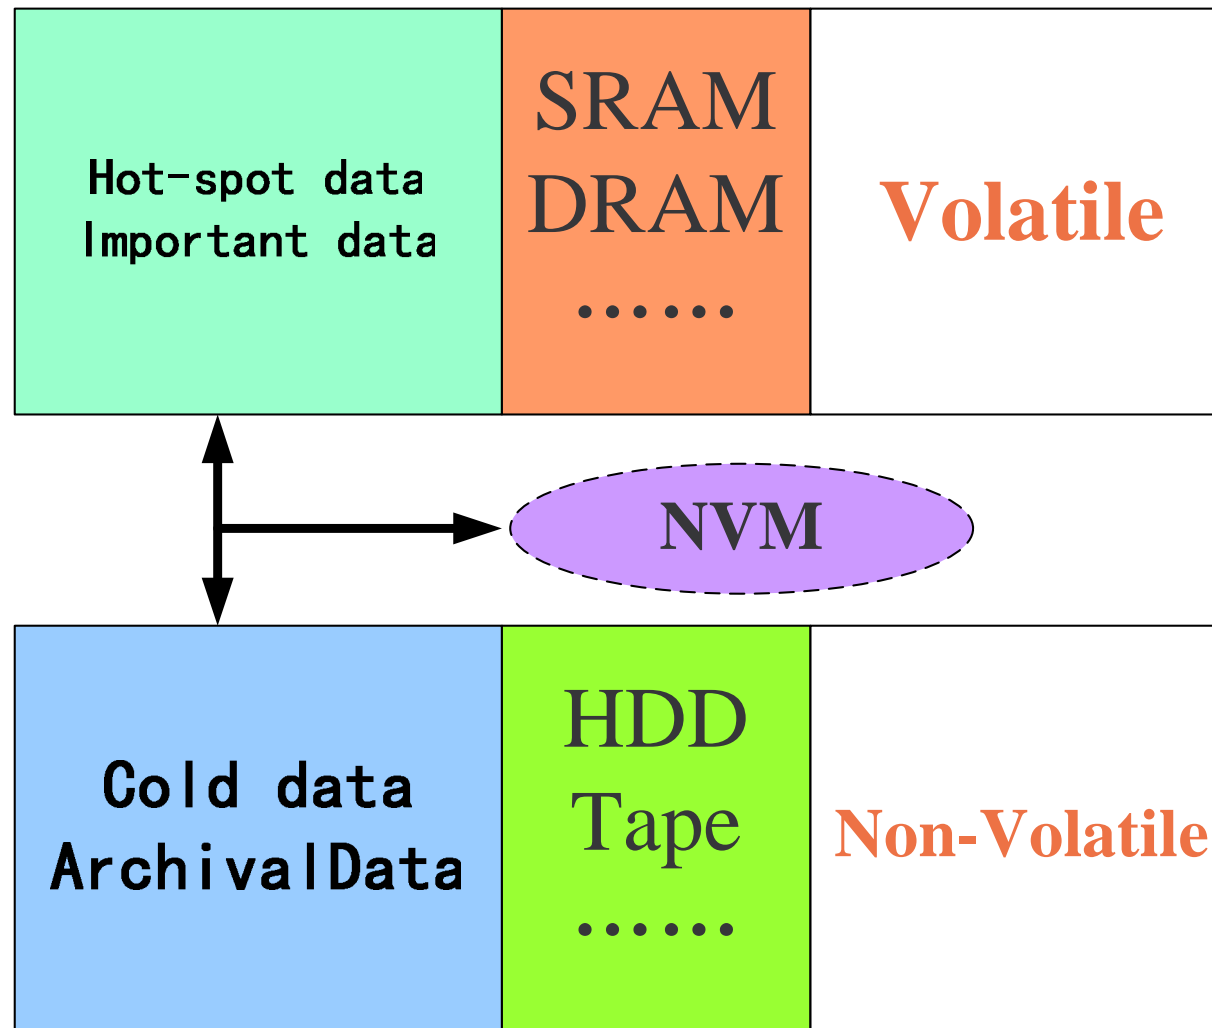
Current and Future Storage 5

- SMILE
- Edge :
- Edge computing, fog computing, proximity computing,


Challenge: Hierarchical Architecture

- Heterogeneous Principle
- Differentiated Performance
- Management Complexity
- One-storey house->Skyscraper
- More and more levels


Challenge: Storage Reliability



Hierarchical Data Structure



This tree is too **FAT** !




This tree is too **HIGH** !

Hierarchical and Vertical Architecture

- **Idea:** based on locality principle, some key data consume many system resources.
- However, in the era of big data, the efficiency of locality becomes weak, thus being difficult to improve hit ratio.

The essence behind Hierarchy

- Goal : identify the correlation
- In essence, the hierarchy is an approach to dynamic filter data to obtain correlated aggregation and on-demand allocation.
- If the flat or semi-hierarchical schemes are able to achieve the same goal, it would be much better with significant performance improvements.



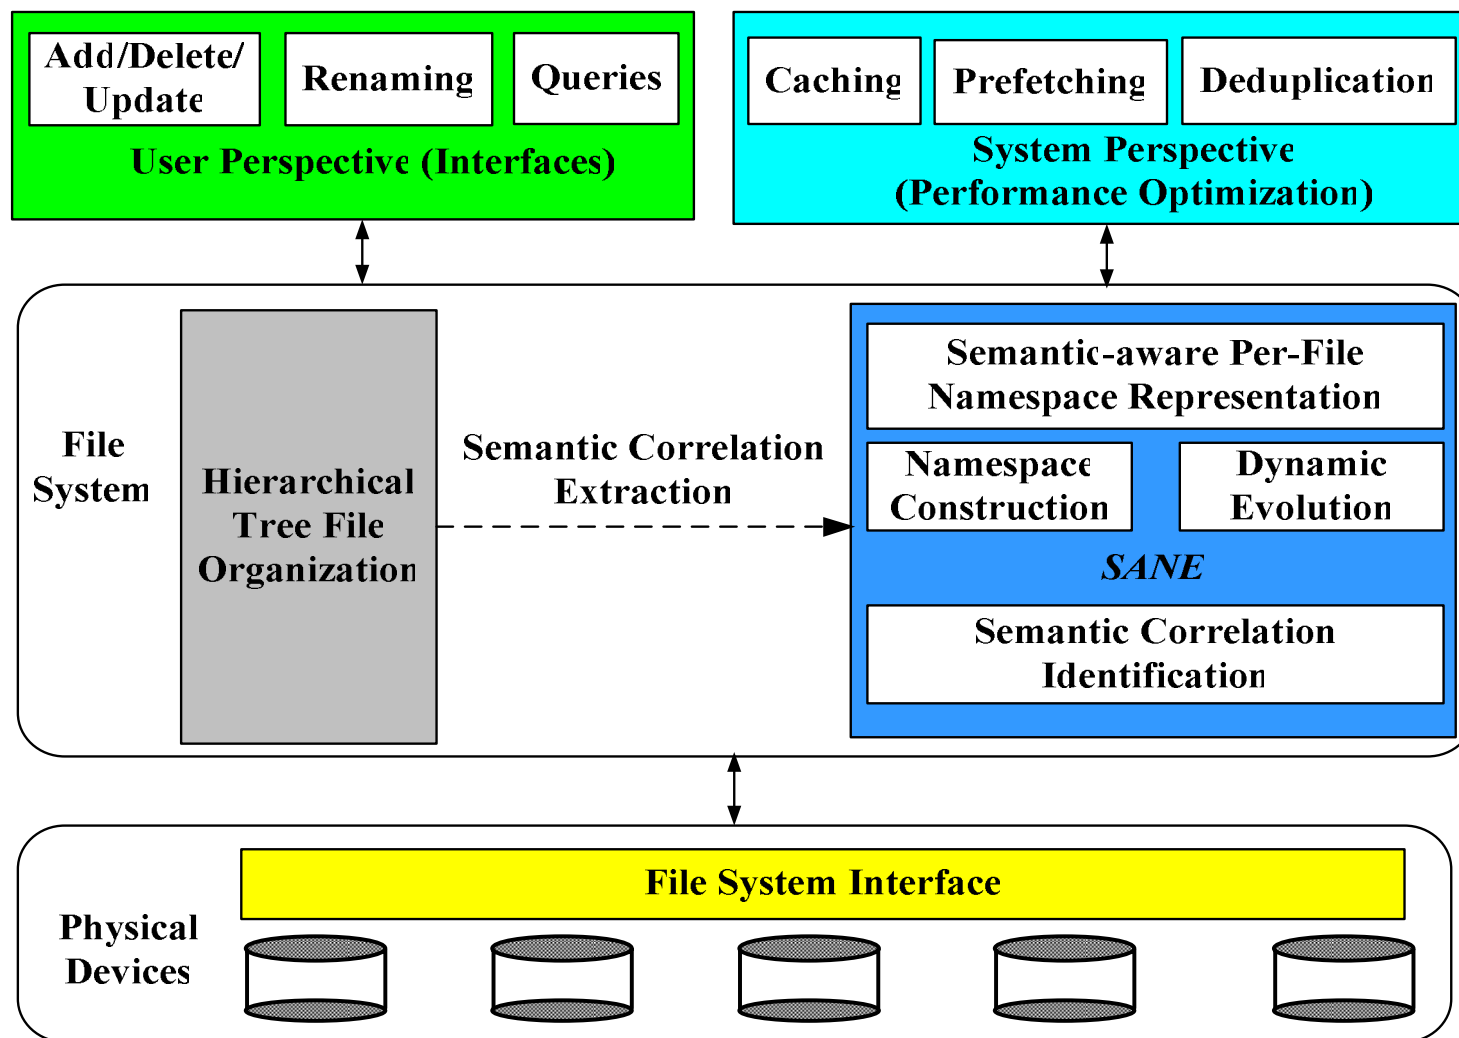
Semi-hierarchical Architecture

- Problem to be addressed:
 - How to storage data in large-scale storage systems
- The idea:
 - Semantic storage is the new form of implementing storage systems.

Our related work

- **Semantic Namespace** : SANE(TPDS14)
- **Semantic Aggregation** : FAST(SC14), HAR(ATC14), SiLo(ATC11),
- **Semantic Hash Computation** : SmartCuckoo(ATC17), DLSH(SoCC17), SmartEye(INFOCOM15), NEST(INFOCOM13)
- **Semantic On-line Service** : ANTELOPE(TC14)

SANE: The namespace

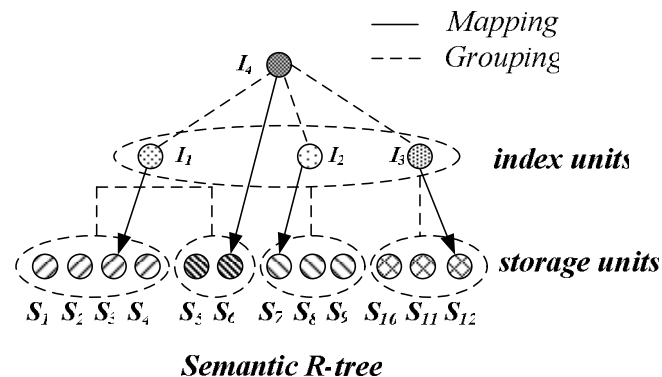

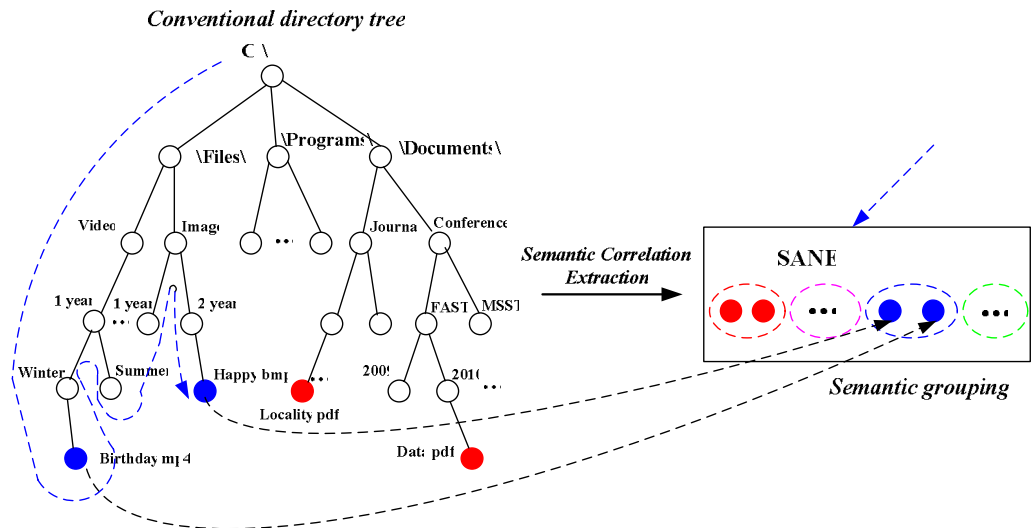


"SANE: Semantic-Aware Namespace in Ultra-large-scale File Systems", *IEEE Transactions on Parallel and Distributed Systems (TPDS)*, Vol.25, No.5, May 2014, pages:1328-1338.

SANE: The Semantic Namespace



Flat Addressing

- Hierarchy becomes the performance bottleneck
- **Design goals :**
 - Searchable
 - Unique

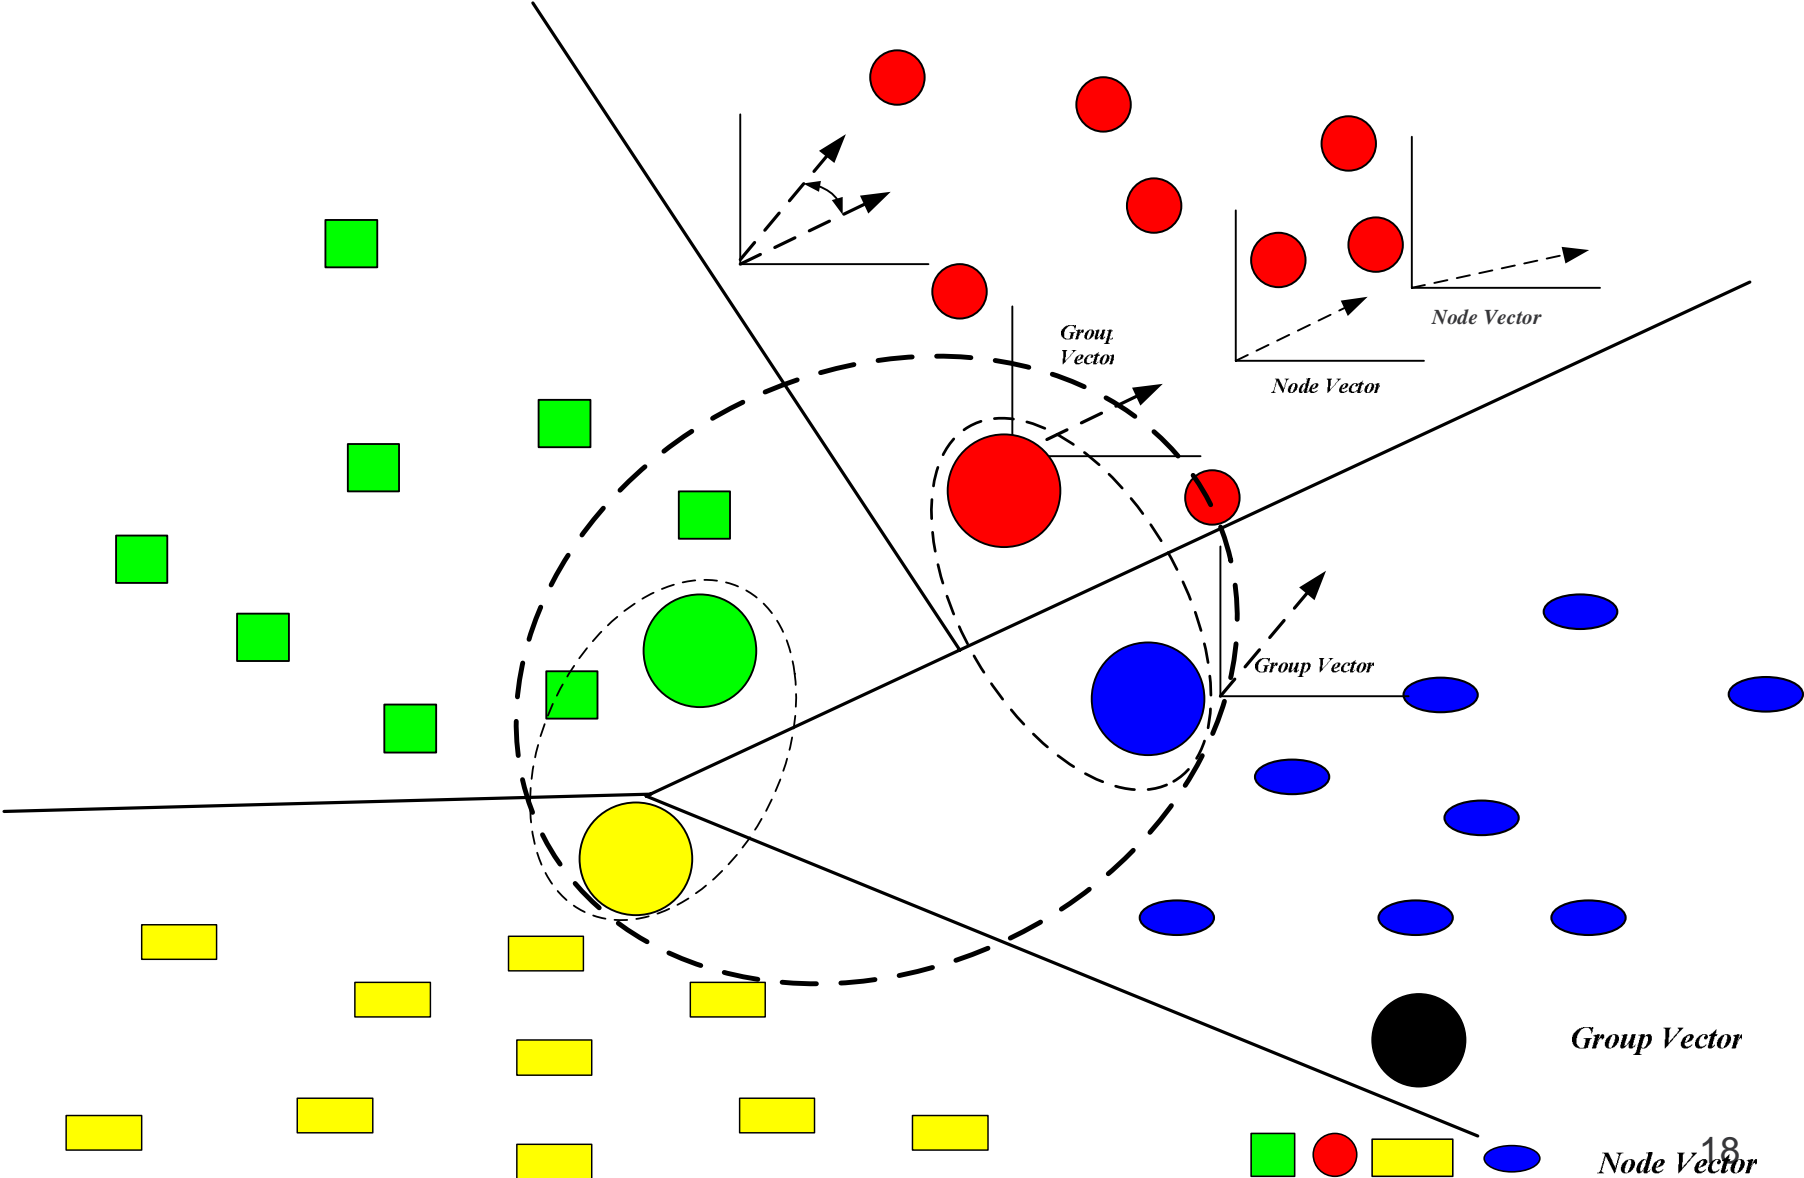


- Construct the semantic-aware namespace

Comparisons with Conventional File Systems

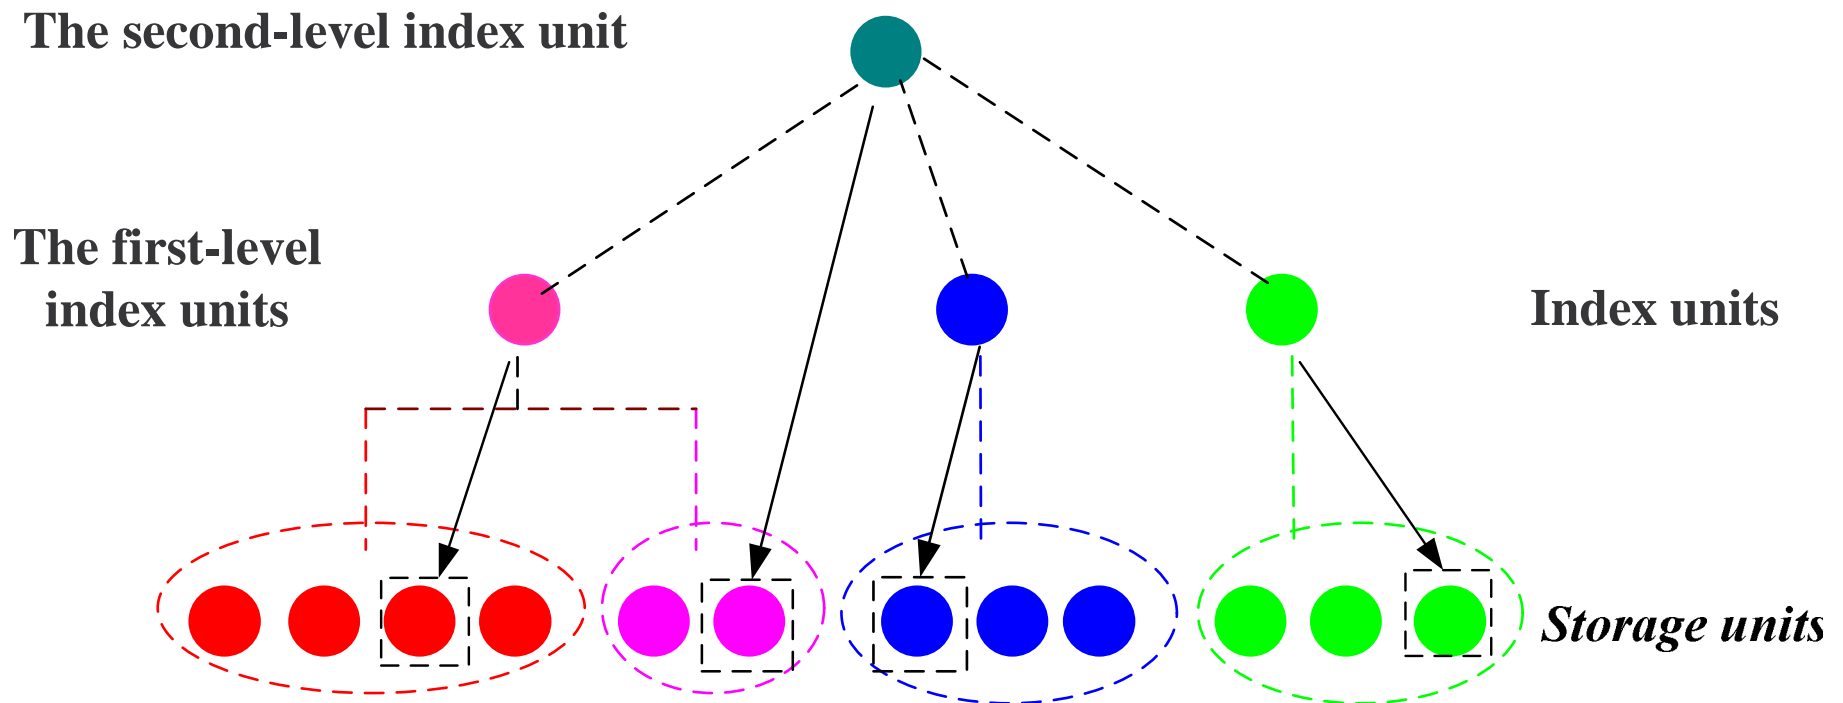


Grouping Procedures

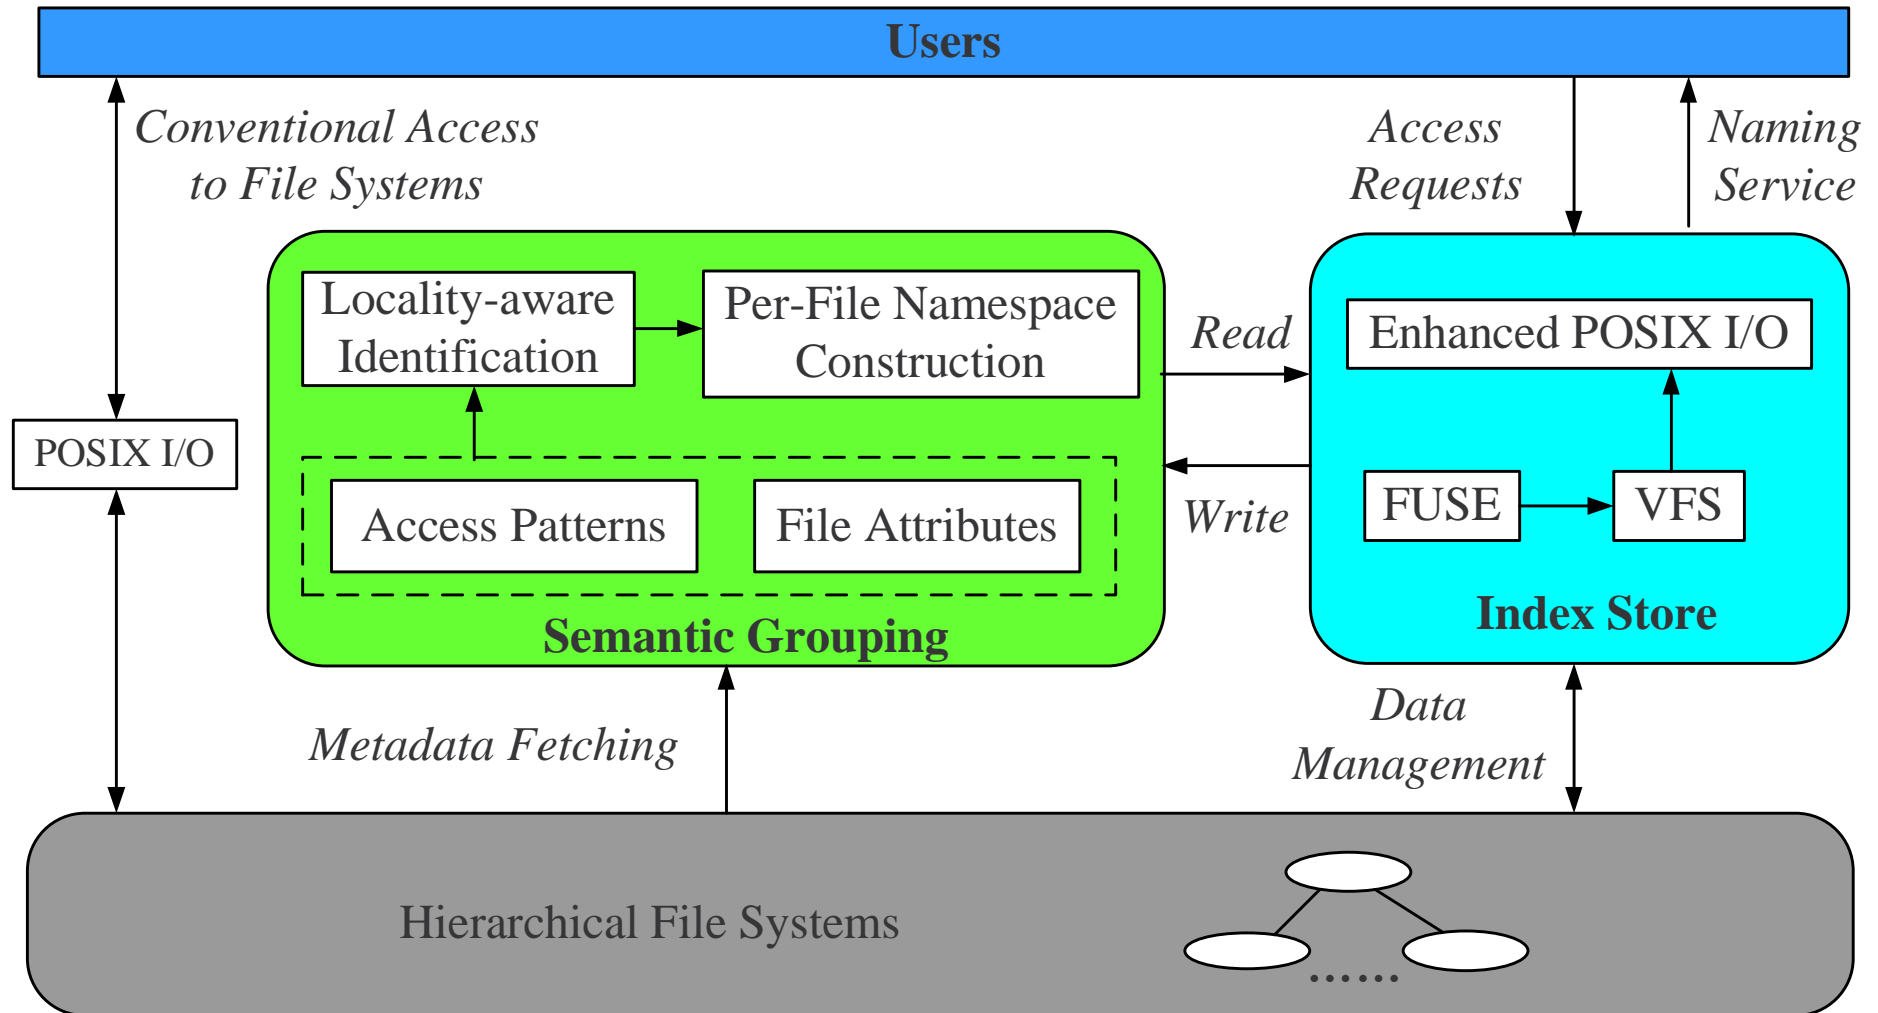


Mapping of Index Units

- **Our mapping is based on a simple bottom-up approach that iteratively applies random selection and labeling operations.**




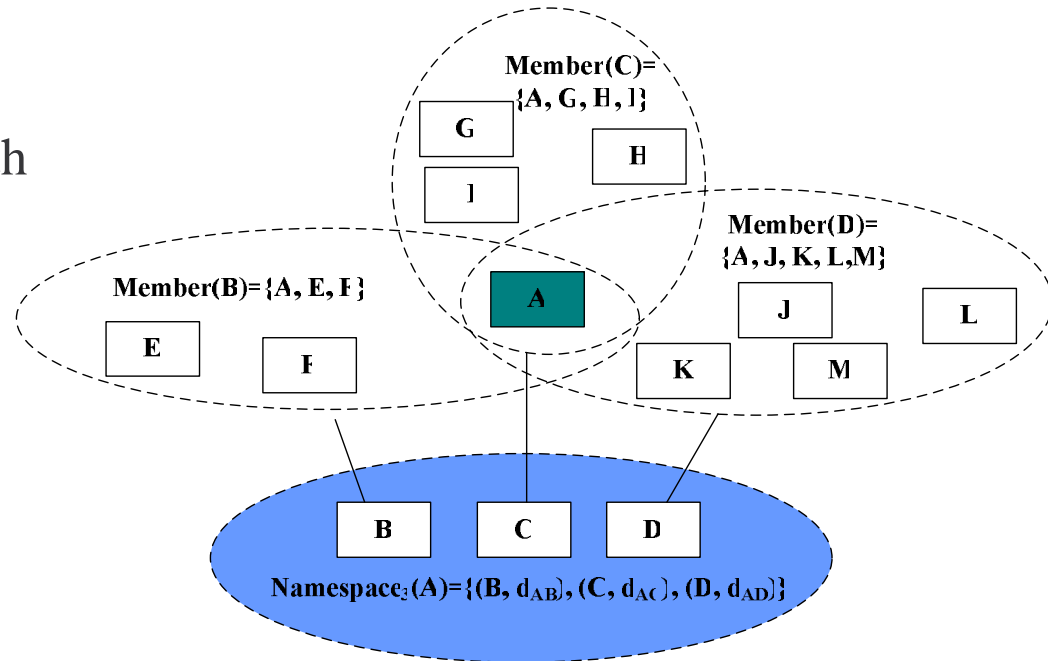
Components



Naming and Rename

- **Submodular Maximization**
- Select a subset of namespaces with distinct names

$$S^* \in \operatorname{argmax}_{S \subseteq V} F(S) \quad s.t. |S| \leq |T|.$$



Namespace(A) = {(E, d_{AB}), (C, d_{AC}), (D, d_{AE})}

Member(B) = {(A, d_{AB}), (E, d_{BE}), (F, d_{BF})}

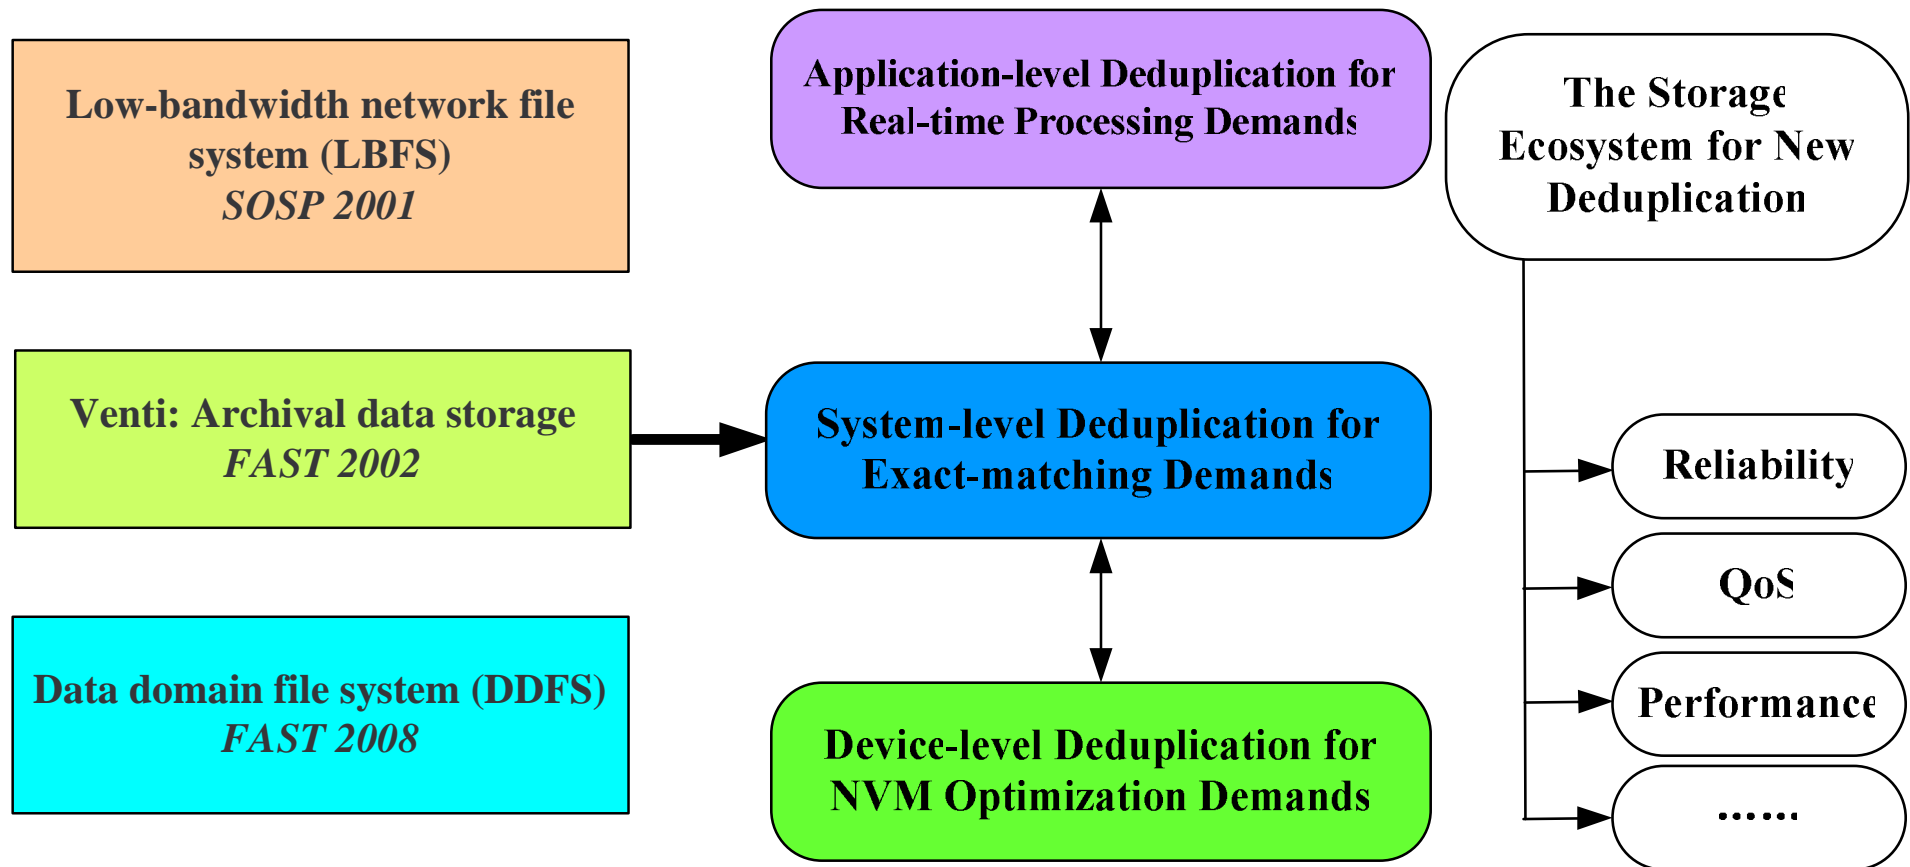
Member(C) = {(A, d_{AC}), (G, d_{CG}), (H, d_{CH}), (I, d_{CI})}

Member(D) = {(A, d_{AD}), }


Maximization for Monotone Submodular functions

- Scoring Function is a monotone submodular function
 - Greedy algorithm
 - Constant-scale mathematical quality guarantee

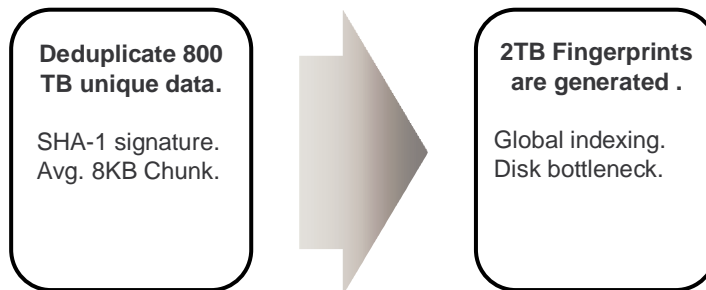
Example 1: New Deduplication Ecosystem




The Synergization of Similarity and Locality—SiLo



The Scalability of Deduplication Indexing



- Expose and exploit more similarity by grouping strongly correlated small files into a segment and segmenting large files
- Leverage locality in the backup stream by grouping contiguous segments into blocks to capture similar and duplicate data missed by the probabilistic similarity detection.



	Small files ($\leq 64\text{KB}$)	Large files ($\geq 2\text{MB}$)
Percentage of total file number	$\geq 80\%$	$\leq 20\%$
Percentage of total space	$\leq 20\%$	$\geq 80\%$

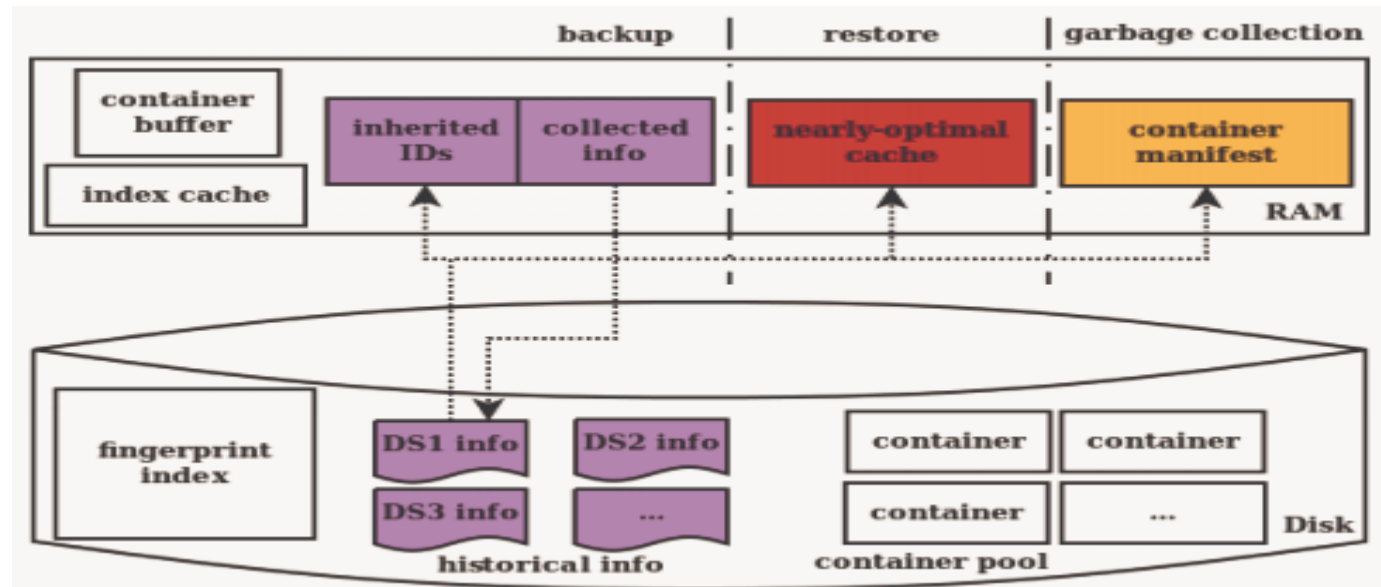
Grouping many highly correlated small files into a segment to minimize dedup overheads

Dividing the large files into many small segments to expose more similarity characteristics

“SiLo: A Similarity-Locality based Near-Exact Deduplication Scheme with Low RAM Overhead and High Throughput,” Proceedings of USENIX ATC, June 2011.

Fragmentation in Deduplication

- History-Aware Rewriting algorithm (HAR)

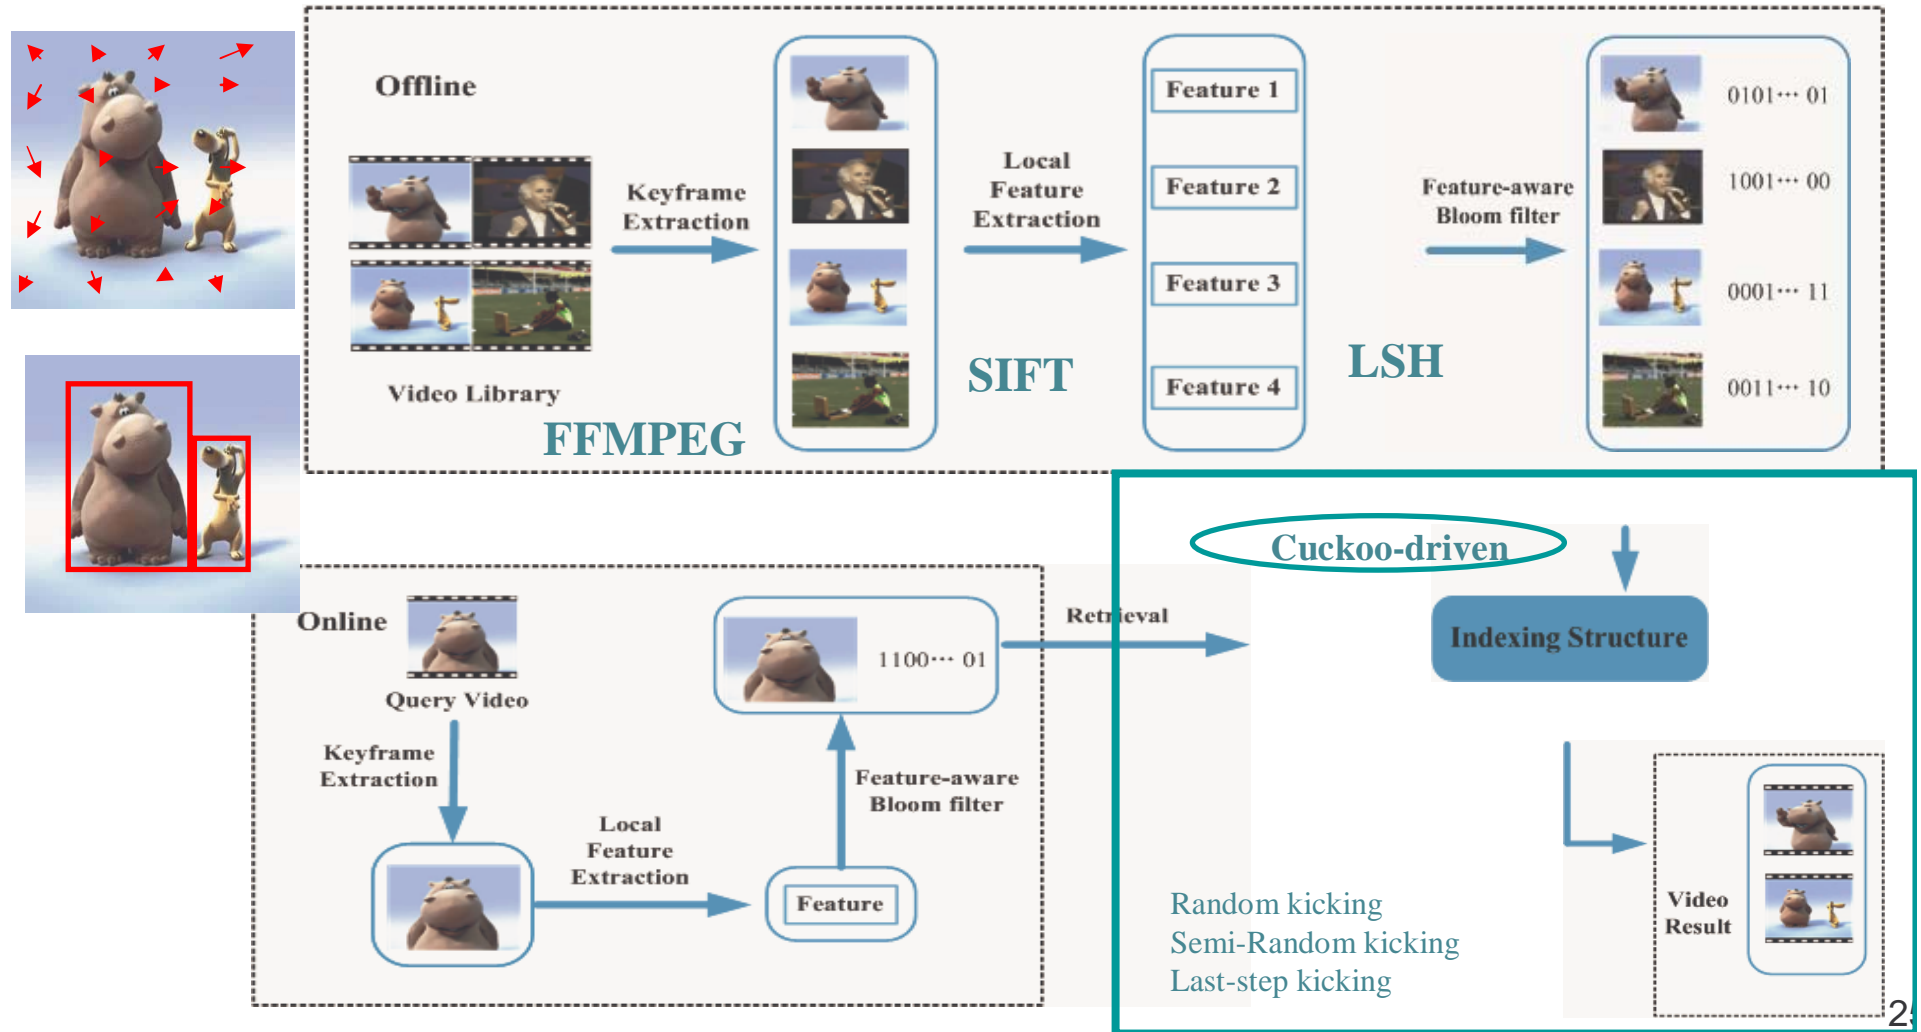


- The fragmentation decreases restore performance and results in invalid chunks becoming physically scattered in different containers after users delete backups.
- HAR exploits historical information of backup systems to more accurately identify and rewrite fragmented chunks.

24

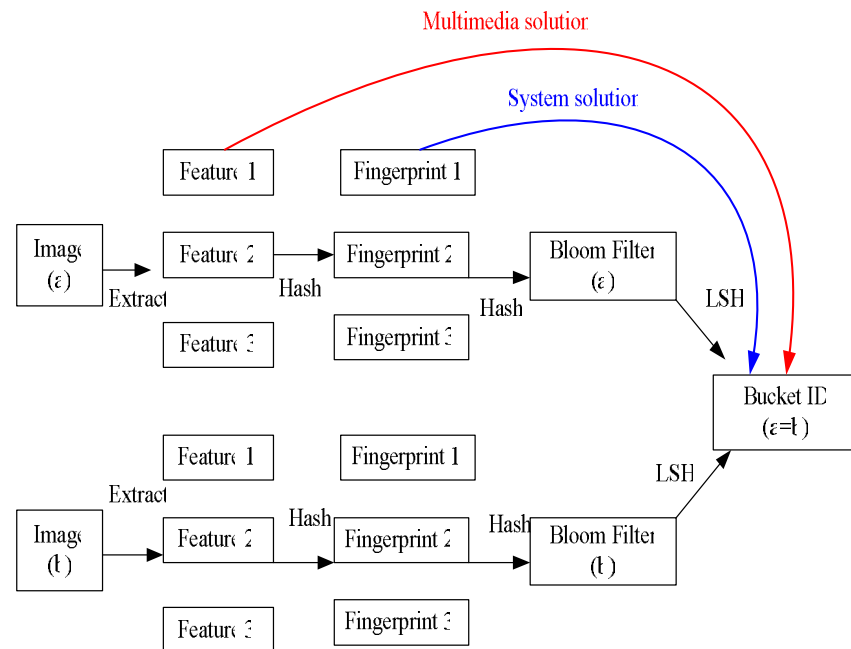

"Accelerating Restore and Garbage Collection in Deduplication-based Backup Systems via Exploiting Historical Information", Proc. USENIX ATC, 2014,

Example 2: Application-level Approximate Methodology--FAST



"FAST: Near Real-time Searchable Data Analytics for the Cloud", Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis (SC), November 2014


Application-level Approximate Methodology: FAST




"FAST: Near Real-time Searchable Data Analytics for the Cloud", Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis (SC), November 2014

Approximate Image Transmission in Networking: SmartEye


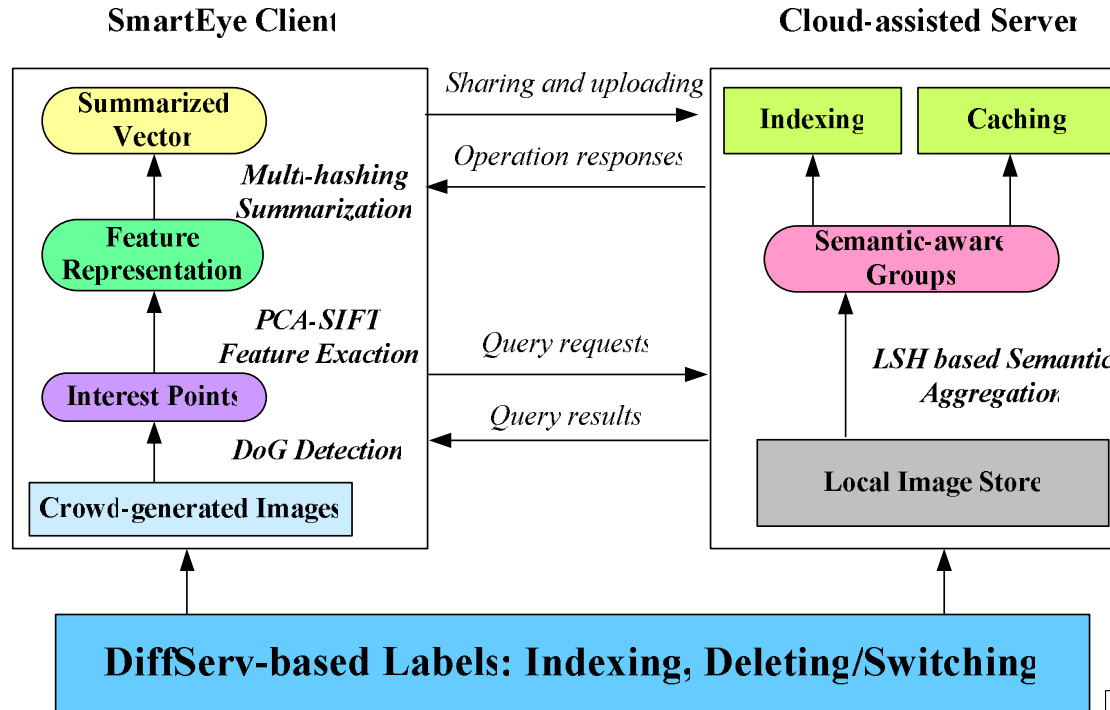

Similarity



Label



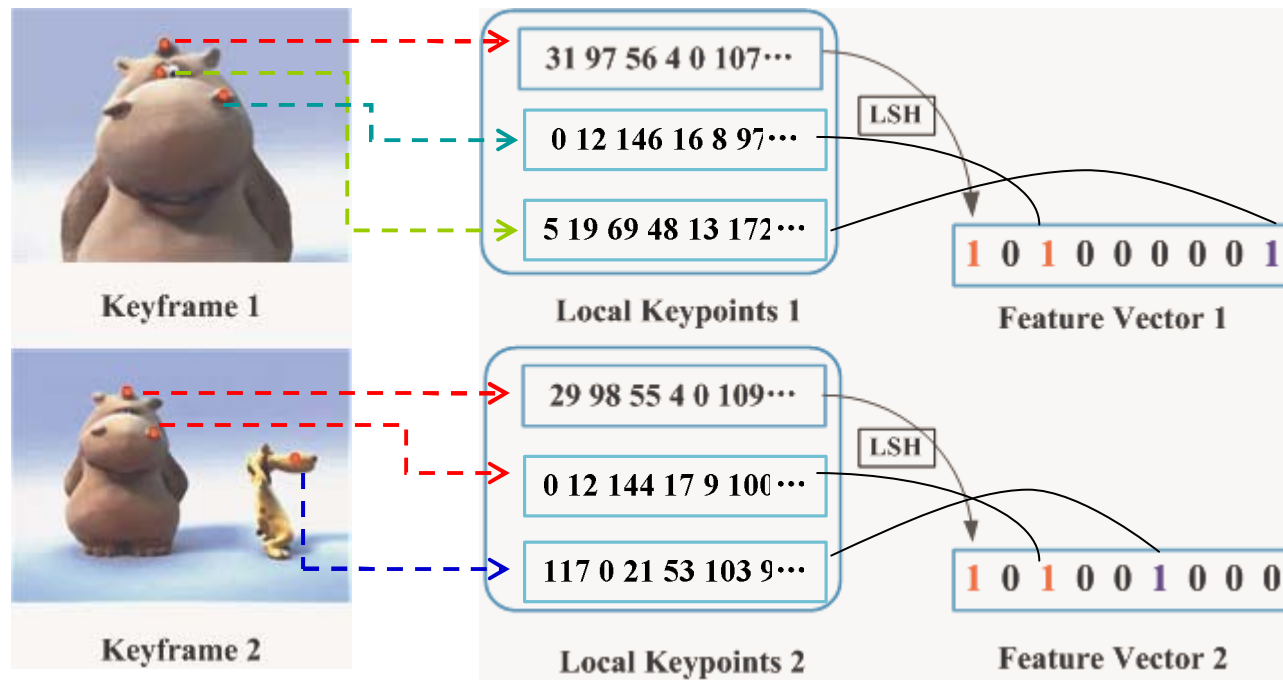
Operations



"SmartEye: Real-time and Efficient Cloud Image Sharing for Disaster Environments" 27
 Proceedings of INFOCOM, 2015, pages: 1616-1624

The design of SmartEye

- Compact Feature Representation




Locality Sensitive Hashing (LSH)


- If $\|p, q\|_s \leq R$ then $Pr_{\mathbb{H}}[h(p) = h(q)] \geq P_1$,
- If $\|p, q\|_s > cR$ then $Pr_{\mathbb{H}}[h(p) = h(q)] \leq P_2$.

Near neighbor?

- *yes*
- *not sure*
- *no*



Locality-Sensitive Hashing (LSH)




- Close items will collide with high probability
- Distant items will have very little chance to collide

Efficient Cuckoo-driven LSH


- (1) Use Cuckoo Driven LSH to reduce search time when collision occurs
- (2) Use neighbor buckets to further reduce the possibility of kickout
- (3) Space efficiency due to neighboring probe and data locality

Blue: hit position by LSH computation
Green: Neighbor bucket has data correlation



A multi-choice LSH

If all $LSH_i(a)$ are full, can choose adjacent empty bucket



Available locations for item a

Probing adjacent neighbors: the probability of endless “kicking out” is much more smaller than ordinary cuckoo hashing


NEST: Efficient Cuckoo-driven LSH

- (1) Use Cuckoo Driven LSH to reduce search time when collision occurs
- (2) Use neighbor buckets to further reduce the possibility of kickout
- (3) Space efficiency due to neighboring probe and data locality


Blue: hit position by LSH computation

Green: Neighbor bucket has data correlation

If all $LSH_i(a)$ are full, can choose adjacent empty bucket




A multi-choice LSH




Available locations for item a

Probing adjacent neighbors: the probability of endless “kicking out” in NEST is much more smaller than ordinary cuckoo hashing

NEST: Resolve Collision: if still fails




Hashing collisions for inserting item *a*




Moving item *h* to its another location

Note: Adjacent probing significantly reduce or even avoid hash failing (FAST INDEX)




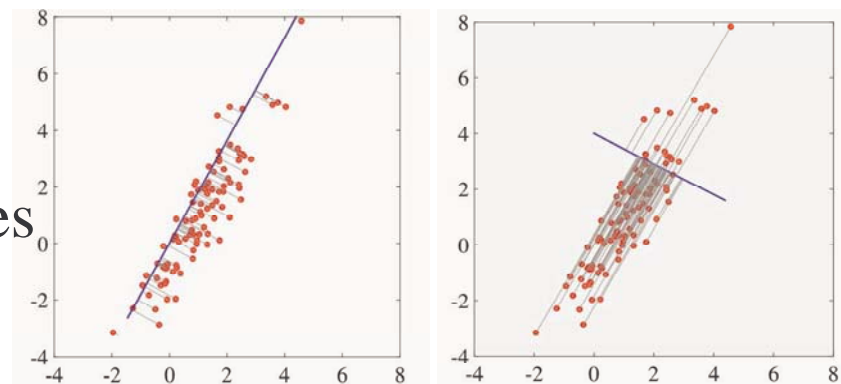
Example 3: Pseudoforest



- An **endless loop** is formed.
- **Endless kickouts** for any insertion within the loop.

DLSH: A Distribution-aware LSH

- Due to distribution-unaware projection vectors:
 - Multiple hash tables to maintain data locality and guarantee the query accuracy.
- Design goal:
 - Decrease the number of hash tables
 - Mitigate in-memory consumption
- Approach:
 - Differentiating the aggregated data in a suitable direction;
 - Exhibiting the data locality as well as decreasing the hash collisions.



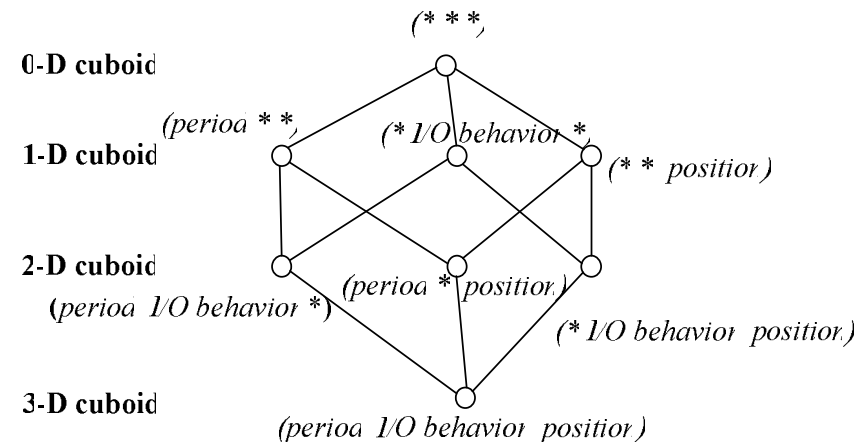
"DLSH: A Distribution-aware LSH Scheme for Approximate Nearest Neighbor Query in Cloud Computing", *Proceedings of ACM Symposium on Cloud Computing (SoCC)*, 2017

Example 4: On-line Precomputation--Data Cube

Time	Evening	10	22	6	8	Direction
	Afternoon	57	196	188	261	
	Morning					
Direction	West	15	176	168	52	5
	East	18	158	172	69	8
	South	56	20	127	82	9
	North	28	372	165	55	
		Str A	Str E	Str C	Str D	
Position						

Period	<i>Evening</i>	10	22	6	8
	<i>Afternoon</i>	57	196	188	152
	<i>Morning</i>				
I/O Behavior	<i>Read</i>	56	206	127	82
	<i>Write</i>	28	372	165	55
		<i>Server A</i>	<i>Server E</i>	<i>Server C</i>	<i>Server D</i>
		Position			

- Leverage precomputation based data cube to support online cloud services
- Use semantic-aware partial materialization to reduce the operation and space overheads



"ANTELOPE: A Semantic-aware Data Cube Scheme for Cloud Data Center Networks", 36
 IEEE Transactions on Computers (TC), Vol.63, No.9, September 2014, pages: 2146-2159.

Open Source Codes (in GitHub)

- [SmartCuckoo](#): in GitHub. *SmartCuckoo is a new cuckoo hashing scheme to support metadata query service.*
 - <https://github.com/syy804123097/SmartCuckoo>
- [SmartSA](#) (E-STORE): in GitHub to support near-deduplication for image sharing based on the energy availability in Smartphone.
 - <https://github.com/Pfzuo/SmartSA>
- [Real-time-Share](#): in GitHub, to support real-time image sharing in the cloud, which is an important component of [SmartEye](#) (INFOCOM 2015).
 - <https://github.com/syy804123097/Real-time-Share>
- [MinCounter](#): in GitHub. *MinCounter is the proposed data structure in the [MSST 2015 Paper](#).*
 - <https://github.com/syy804123097/MinCounter>
- [NEST](#): in GitHub (Download [INFOCOM 2013 Paper](#), [Source Codes](#), [Manual](#) and [TraceData](#)).
 - <https://github.com/syy804123097/NEST>
- [LSBF](#) (Locality-Sensitive Bloom Filter): in GitHub (Download [TC 2012 Paper](#), [Source Codes](#) and [Manual](#)).
 - <https://github.com/syy804123097/LSBF>

Thanks and Questions